

# Linear Programming and Game Theory

Ron Parr  
CPS 570


With thanks to Vince Conitzer for some content

## What are Linear Programs?

- Linear programs are **constrained optimization problems**
- Constrained optimization problems ask us to maximize or minimize a function subject to mathematical constraints on the variables
  - Convex programs have convex objective functions and convex constraints
  - Linear programs (special case of convex programs) have linear objective functions and linear constraints
- LPs = generic language for wide range problems
- LP solvers = widely available hammers
- Entire classes and vast expertise invested in making problems look like nails

## Linear programs: example

- Make reproductions of 2 paintings



- Painting 1:
  - Sells for \$30
  - Requires 4 units of blue, 1 green, 1 red
- Painting 2
  - Sells for \$20
  - Requires 2 blue, 2 green, 1 red
- We have 16 units blue, 8 green, 5 red

$$\text{maximize } 3x + 2y$$

subject to

$$4x + 2y \leq 16$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$$x \geq 0$$

$$y \geq 0$$

## Solving the linear program graphically

$$\text{maximize } 3x + 2y$$

subject to

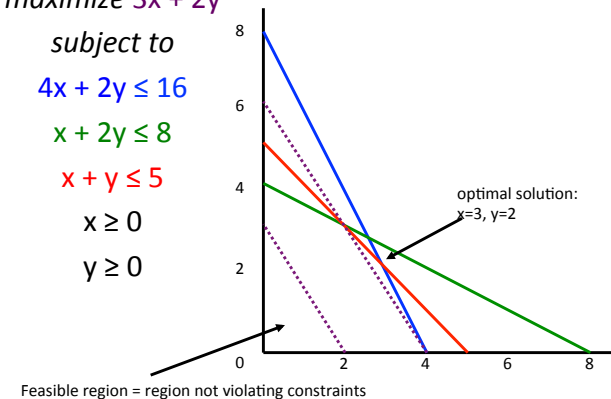
$$4x + 2y \leq 16$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$$x \geq 0$$


$$y \geq 0$$



## Linear Programs in General

- Linear constraints, linear objective function
  - Maximize (minimize):  $f(\mathbf{x})$  ← Linear function of vector  $\mathbf{x}$
  - Subject to:  $\mathbf{Ax} \leq \mathbf{b}$   
     ← Matrix  $\mathbf{A}$
- Can swap maximize/minimize,  $\leq/\geq$ ; can add equality
- View as search: Searches space of values of  $\mathbf{x}$
- Alternatively: Search for tight constraints w/high objective function value

## What Happens In Higher Dimensions (1) Understanding the Feasible Region



## What Happens In Higher Dimensions (2) lines → hyperplanes

- Inequality w/2 variables → one side of a line
- 3 variables → one side of a plane
- k variables → one side of hyperplane
- Physical intuition:



<http://www.rubylane.com/Item/623546-4085/Orrefors-x22Zenithx22-Pattern-Crystal-Bowl>

## Solving linear programs (1)

- Optimal solutions always exist at vertices of the feasible region
  - Why?
  - Assume you are not at a vertex, you can always push further in direction that improves objective function (or at least doesn't hurt)
  - How many vertices does a  $k \times n$  matrix imply?
- Dumb(est) algorithm:
  - Given  $n$  variables,  $k$  constraints
  - Check all  $k$ -choose- $n = O(k^n)$  possible vertices

### Solving linear programs (2)

- Smarter algorithm (simplex)
  - Pick a vertex
  - Repeatedly hop to neighboring (one different tight constrain) vertices that improve the objective function
  - Guaranteed to find solution (no local optima)
  - May take exponential time in worst case (though rarely)
- Still smarter algorithm
  - Move inside the interior of the feasible region, in direction that increases objective function
  - Stop when no further improvements possible
  - Tricky to get the details right, but weakly polynomial time

### Solving LPs in Practice

- Use commercial products like cplex or gurobi
- Do not try to implement an LP solver yourself!
- Do not use matlab's linprog for anything other than small problems. Really. No – REALLY!

### Modified LP


$\text{maximize } 3x + 2y$   
 $\text{subject to}$   
 $4x + 2y \leq 15$   
 $x + 2y \leq 8$   
 $x + y \leq 5$   
 $x \geq 0$   
 $y \geq 0$

Optimal solution:  $x = 2.5, y = 2.5$   
 Solution value =  $7.5 + 5 = 12.5$   
 Half paintings?

### Integer (linear) program

$\text{maximize } 3x + 2y$   
 $\text{subject to}$   
 $4x + 2y \leq 15$   
 $x + 2y \leq 8$   
 $x + y \leq 5$   
 $x \geq 0, \text{ integer}$   
 $y \geq 0, \text{ integer}$

optimal IP solution:  $x=2, y=3$  (objective 12)  
 optimal LP solution:  $x=2.5, y=2.5$  (objective 12.5)



## Mixed integer (linear) program

maximize  $3x + 2y$

subject to


$$4x + 2y \leq 15$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$$x \geq 0$$

$$y \geq 0, \text{ integer}$$



## Solving (M)IPs

- (Mixed) Integer programs are NP-hard to solve
- Intuition: Constraint surface is jagged; no obvious way to avoid checking exponential number of assignments to integer variables
- In practice:
  - Constraints often give clues on how to restrict number of solutions considered
  - Smart solvers (cplex, gurobi) can *sometimes* find solutions to large (M)IPs surprisingly quickly (and surprisingly slowly)

## LP Trick (one of many)

- Suppose you have a huge number of constraints, but a small number of variables ( $k \gg n$ )
- Constraint generation:
  - Start with a subset of the constraints
  - Find solution to simplified LP
  - Find most violated constraint, add back to LP
  - Repeat
- Why does this work?
  - If missing constraints are unviolated, then adding them back wouldn't change the solution
  - Sometimes terminates after adding in only a fraction of total constraints
  - No guarantees, but often helpful in practice

## Duality

- For every LP there is an equivalent “Dual” problem
- Solution to primal can be used to reconstruct solution to dual, and vice versa
- LP duality:

$$\begin{array}{ll}
 \text{minimize : } c^T x & \longleftrightarrow \text{maximize : } b^T y \\
 \text{subject to : } Ax = b & \text{subject to : } A^T y = c \\
 : x \geq 0 & : y \geq 0
 \end{array}$$

## MDP Solved as an LP

$$V(s) = \max_a R(s,a) + \gamma \sum_{s'} P(s'|s,a) V(s')$$

Issue: Turn the non-linear max into a collection of linear constraints

$$\forall s,a : V(s) \geq R(s,a) + \gamma \sum_{s'} P(s'|s,a) V(s')$$

MINIMIZE:  $\sum_s V(s)$

Optimal action has  
tight constraints

## What is Game Theory? I

- Very general mathematical framework to study situations where multiple agents interact, including:
  - Popular notions of games
  - Everything up to and including multistep, multiagent, simultaneous move, partial information games
  - Example Duke CS research: Aiming sensors to catch hiding enemies, assigning guards to posts
  - Can even include negotiating, posturing and uncertainty about the players and game itself

- von Neumann and Morgenstern (1944) was a major launching point for modern game theory
- Nash: Existence of equilibria in [general sum](#) games



(wikipedia)

## What is game theory? II

- Study of settings where multiple agents each have
  - Different preferences (utility functions),
  - Different actions
- Each agent's utility (potentially) depends on all agents' actions
  - What is optimal for one agent depends on what other agents do
  - Can be circular
- Game theory studies how agents can rationally form [beliefs](#) over what other agents will do, and (hence) how agents should [act](#)
- Useful for acting and (potentially) predicting behavior of others
- Not necessarily descriptive

## Real World Game Theory Examples

- War
- Auctions
- Animal behavior
- Networking protocols
- Peer to peer networking behavior
- Road traffic
- Mechanism design:
  - Suppose we want people to do X?
  - How to engineer situation so they will act that way?

## Covered Today

- 2 player, zero sum **simultaneous move** games
- Example: Rock, Paper, Scissors
- Linear programming solution







## Linear Programs (max formulation)

$$\begin{aligned} &\text{maximize : } c^T x \\ &\text{subject to : } Ax \leq b \\ &\quad \quad \quad : x \geq 0 \end{aligned}$$

- Note: min formulation also possible
  - Min:  $c^T x$
  - Subject to:  $Ax \geq b$
- Some use equality as the canonical representation (introducing slack variables)
- LP tricks
  - Multiply by -1 to reverse inequalities
  - Can easily introduce equality constraints, or arbitrary domain constraints

## Rock, Paper, Scissors Zero Sum Formulation

- In zero sum games, one player's loss is other's gain
- Payoff matrix:

				
		<i>R</i>	<i>P</i>	<i>S</i>
	<i>R</i>	0	-1	1
	<i>P</i>	1	0	-1
	<i>S</i>	-1	1	0

- Minimax solution maximizes worst case outcome

## Rock, Paper, Scissors Equations

- $R, P, S$  = probability that we play rock, paper, or scissors respectively ( $R + P + S = 1$ )
- $U$  is our expected utility
- Bounding our utility:
  - Opponent rock case:  $U \leq P - S$
  - Opponent paper case:  $U \leq S - R$
  - Opponent scissors case:  $U \leq R - P$
- Want to maximize  $U$  subject to constraints
- Solution:  $(1/3, 1/3, 1/3)$

### Rock, Paper, Scissors LP Formulation

- Our variables are:  $x=[U,R,P,S]^T$
- We want:
  - Maximize  $U$
  - $U \leq P - S$
  - $U \leq S - R$
  - $U \leq R - P$
  - $R+P+S = 1$
- How do we make this fit:  $\boxed{\begin{array}{l} \text{maximize : } c^T x \\ \text{subject to : } Ax \leq b \\ \phantom{\text{subject to : }} : x \geq 0 \end{array}} \quad ?$

### Rock Paper Scissors LP Formulation

$$x = [U, R, P, S]^T$$

$$A = \begin{pmatrix} 1 & 0 & -1 & 1 \\ 1 & 1 & 0 & -1 \\ 1 & -1 & 1 & 0 \\ 0 & 1 & 1 & 1 \\ 0 & -1 & -1 & -1 \end{pmatrix}$$

$$b = [0, 0, 0, 1, -1]^T$$

$$c = [1, 0, 0, 0]^T$$

$$\boxed{\begin{array}{l} \text{maximize : } c^T x \\ \text{subject to : } Ax \leq b \\ \phantom{\text{subject to : }} : x \geq 0 \end{array}}$$

### Rock, Paper, Scissors Solution

- If we feed this LP to an LP solver we get:
  - $R=P=S=1/3$
  - $U=0$
- Solution for the other player is:
  - The same...
  - By symmetry
- This is the minimax solution
- This is also an equilibrium
  - No player has an incentive to deviate
  - (Defined more precisely later)

### Tangent: Why is RPS Fun?

- OK, it's not...
- Why *might* RPS be fun?
  - Try to exploit non-randomness in your friends
  - Try to be random yourself

## Minimax Solutions in General


- What do we know about minimax solutions?
  - Can a suboptimal opponent trick minimax?
  - When should we abandon minimax?
- Minimax solutions for 2-player zero-sum games can always be found by solving a linear program
- The minimax solutions will also be equilibria
- For general sum games:
  - Minimax does not apply
  - Equilibria may not be unique
  - Need to search for equilibria using more computationally intensive methods

## Outline


- Digression: Linear Programming
- 2 player, zero sum **simultaneous move** games
- Example: Rock, Paper, Scissors
- Linear programming solution
- General sum games

## “Chicken”

- Two players drive cars towards each other
- If one player goes straight, that player wins
- If both go straight, they both die



## Rock-paper-scissors – Seinfeld variant



MICKEY: All right, rock beats paper!  
 (Mickey smacks Kramer's hand for losing)  
 KRAMER: I thought paper covered rock.  
 MICKEY: Nah, rock flies right through paper.  
 KRAMER: What beats rock?  
 MICKEY: (looks at hand) Nothing beats rock.

	Rock	Paper	Scissors
Rock	0, 0	1, -1	1, -1
Paper	-1, 1	0, 0	-1, 1
Scissors	-1, 1	1, -1	0, 0



## Dominance

- Player  $i$ 's strategy  $s_i$  **strictly dominates**  $s_i'$  if
  - for any  $s_{-i}$ ,  $u_i(s_i, s_{-i}) > u_i(s_i', s_{-i})$
- $s_i$  **weakly dominates**  $s_i'$  if
  - for any  $s_{-i}$ ,  $u_i(s_i, s_{-i}) \geq u_i(s_i', s_{-i})$ ; and
  - for some  $s_{-i}$ ,  $u_i(s_i, s_{-i}) > u_i(s_i', s_{-i})$

$-i$  = "the player(s) other than  $i$ "

	0, 0	1, -1	1, -1
	-1, 1	0, 0	-1, 1
	-1, 1	1, -1	0, 0

## Prisoner's Dilemma

- Pair of criminals has been caught
- District attorney has evidence to convict them of a minor crime (1 year in jail); knows that they committed a major crime together (3 years in jail) but cannot prove it
- Offers them a deal:
  - If both confess to the major crime, they each get a 1 year reduction
  - If only one confesses, that one gets 3 years reduction

	confess	don't confess
confess	-2, -2	0, -3
don't confess	-3, 0	-1, -1

## "Should I buy an SUV?"

purchasing + gas cost



cost: 5



cost: 3

accident cost

cost: 5



cost: 5

cost: 8



cost: 2

cost: 5



cost: 5

	-10, -10	-7, -11
	-11, -7	-8, -8

## "2/3 of the average" game

- Everyone writes down a number between 0 and 100
- Person closest to 2/3 of the average wins
- Example:
  - A says 50
  - B says 10
  - C says 90
  - Average(50, 10, 90) = 50
  - 2/3 of average = 33.33
  - A is closest ( $|50 - 33.33| = 16.67$ ), so A wins

## Iterated dominance


- Iterated dominance: remove (strictly/weakly) dominated strategy, repeat
- Iterated strict dominance on Seinfeld's RPS:

	Rock	Paper	Scissors
Rock	0, 0	1, -1	1, -1
Paper	-1, 1	0, 0	-1, 1
Scissors	-1, 1	1, -1	0, 0

 $\rightarrow$ 

	Rock	Scissors
Rock	0, 0	1, -1
Scissors	-1, 1	0, 0

## "2/3 of the average" game revisited



## Mixed strategies

- **Mixed strategy** for player  $i$  = **probability distribution** over player  $i$ 's (pure) strategies
- E.g.  $1/3$  Rock,  $1/3$  Paper,  $1/3$  Scissors
- Example of dominance by a mixed strategy:

$1/2$	3, 0	0, 0
$1/2$	0, 0	3, 0
	1, 0	1, 0

## Best Responses

- Let  $A$  be a matrix of player 1's payoffs
- Let  $\sigma_2$  be a mixed strategy for player 2
- $A\sigma_2$  = vector of expected payoffs for each strategy for player 1
- Highest entry indicates **best response** for player 1
- Any mixture of ties is also BR
- Generalizes to  $>2$  players

0, 0	-1, 1	$\sigma_2$
1, -1	-5, -5	

## Nash equilibrium [Nash 50]




- A vector of strategies (one for each player) = a **strategy profile**
- Strategy profile  $(\sigma_1, \sigma_2, \dots, \sigma_n)$  is a **Nash equilibrium** if each  $\sigma_i$  is a **best response** to  $\sigma_{-i}$ 
  - That is, for any  $i$ , for any  $\sigma'_i$ ,  $u_i(\sigma_i, \sigma_{-i}) \geq u_i(\sigma'_i, \sigma_{-i})$
- Does not say anything about multiple agents changing their strategies at the same time
- In any (finite) game, at least one Nash equilibrium (possibly using mixed strategies) exists [Nash 50]
- (Note - singular: equilibrium, plural: equilibria)

## Equilibrium Strategies vs. Best Responses

- equilibrium strategy  $\rightarrow$  best response?
- best response  $\rightarrow$  equilibrium strategy?
- Consider Rock-Paper-Scissors
  - Is  $(1/3, 1/3, 1/3)$  a best response to  $(1/3, 1/3, 1/3)$ ?
  - Is  $(1, 0, 0)$  a best response to  $(1/3, 1/3, 1/3)$ ?
  - Is  $(1, 0, 0)$  a strategy for any equilibrium?


	0, 0	-1, 1	1, -1
	1, -1	0, 0	-1, 1
	-1, 1	1, -1	0, 0

## Nash equilibria of "chicken"



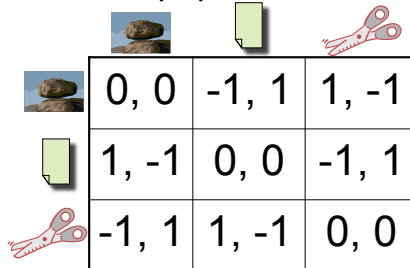
- (D, S) and (S, D) are Nash equilibria
  - They are **pure-strategy Nash equilibria**: nobody randomizes
  - They are also **strict Nash equilibria**: changing your strategy will make you strictly worse off
- No other pure-strategy Nash equilibria







## Equilibrium Selection



- (D, S) and (S, D) are Nash equilibria
- Which do you play?
- What if player 1 assumes (S, D), player 2 assumes (D, S)
- Play is (S, S) = (-5, -5)!!!
- This is the **equilibrium selection** problem

### Rock-paper-scissors



			
	0, 0	-1, 1	1, -1
	1, -1	0, 0	-1, 1
	-1, 1	1, -1	0, 0

- Any pure-strategy Nash equilibria?
- It has a **mixed-strategy Nash equilibrium**:  
Both players put probability 1/3 on each action

### Nash equilibria of "chicken"...

	D	S
D	0, 0	-1, 1
S	1, -1	-5, -5

- Is there a Nash equilibrium that uses mixed strategies -- say, where player 1 uses a mixed strategy?
- If a mixed strategy is a best response, then all of the pure strategies that it randomizes over must also be best responses
- So we need to make player 1 **indifferent** between D and S
- Player 1's utility for playing D =  $-p^c_S$
- Player 1's utility for playing S =  $p^c_D - 5p^c_S = 1 - 6p^c_S$
- So we need  $-p^c_S = 1 - 6p^c_S$  which means  $p^c_S = 1/5$
- Then, player 2 needs to be indifferent as well
- Mixed-strategy Nash equilibrium:  $((4/5 D, 1/5 S), (4/5 D, 1/5 S))$ 
  - People may die! Expected utility -1/5 for each player

$-p^c_S$  = probability  
that column  
player plays s

### Computational Issues

- Zero-sum games - solved efficiently as LP
- General sum games may require exponential time (in # of actions) to find a single equilibrium (no known efficient algorithm and good reasons to suspect that none exists)
- Some better news: Despite bad worst-case complexity, many games can be solved quickly

### Game Theory Issues

- How descriptive is game theory?
  - Some evidence that people play equilibria
  - Also, some evidence that people act irrationally
  - If it is computationally intractable to solve for equilibria of large games, seems unlikely that people are doing this
- How reasonable is (basic) game theory?
  - Are payoffs known?
  - Are situations really simultaneous move with no information about how the other player will act?
  - Are situations really single-shot? (repeated games)
  - How is equilibrium selection handled in practice?

## Extensions

- Partial information
- Uncertainty about the game parameters, e.g., payoffs (Bayesian games)
- Repeated games: Simple learning algorithms can converge to equilibria in some repeated games
- Multistep games with distributions over next states (game theory + MDPs = stochastic games)
- Multistep + partial information (Partially observable stochastic games)
  
- Game theory is so general, that it can encompass essentially all aspects of strategic, multiagent behavior, e.g., negotiating, threats, bluffs, coalitions, bribes, etc.